



Convening stakeholders across industries to craft principles and concrete codes of practice for the development and use of artificial intelligence.

October 11, 2024

RE: Bureau of Industry and Security (BIS) Notice of Proposed Rulemaking (NPRM) on Establishment of Reporting Requirements for the Development of Advanced Artificial Intelligence Models and Computing Clusters

These comments are submitted on behalf of the Alliance for Trust in AI (ATAI), a nonprofit association of companies using artificial intelligence (AI) and representing diverse sectors. Members of ATAI seek to ensure that AI can be a trusted tool by promoting effective policy and clear codes of practice for AI. We appreciate the opportunity to respond to the Bureau of Industry and Security's (BIS) Notice of Proposed Rulemaking (NPRM) on the *Establishment of Reporting Requirements for the Development of Advanced Artificial Intelligence Models and Computing Clusters*.

ATAI appreciates the work that BIS, the Department of Commerce, and the rest of the U.S. government are engaging in to improve the safety and security of AI. BIS has asked for input on the proposed rule regarding reporting requirements for advanced AI models, particularly around the appropriate schedule for reporting and the thresholds for this reporting. ATAI comments on both these elements in the hope that BIS will clarify the broad scope of the proposed rule and re-evaluate the frequency of proposed reporting for covered entities. ATAI also appreciates the need to set a threshold to require reporting on the models that pose the most significant types of risk, but would like to note that these risks cannot be captured by simply describing the model's size. We encourage the development of a more tailored, similarly concrete threshold for this and similar rules.

About ATAI

ATAI brings together companies that use advanced AI in many sectors to advocate for ways that we can build trust in all the kinds of AI that empower companies across the country and the world. ATAI works with companies that are developing foundational AI models, creating AI systems, and implementing these systems and models in their work across industries.

We aim to give organizations concrete guidance on how to build AI responsibly, implement AI principles, support learning and information sharing across sectors, and establish a shared voice for the many users of AI now and in the future. ATAI is building on work done by technologists, policymakers, and academics to create a shared understanding of how to develop and use AI responsibly. Through multi-stakeholder partnerships with members across industries and sectors, ATAI is developing definitions, principles, and codes of practice that ensure that AI is available and trusted for everyone.

Clarity of Scope and Applicability

BIS should make it clear that the proposed rule only applies to the original developers of large models

BIS has requested feedback on the collection thresholds in the proposed rule. In discussing BIS's proposed rule with our members, there have been several points of confusion that we urge BIS to address regarding these thresholds and which organizations might be required to comply. These points include disconnects between the narrative explanation of the rule and the proposed rule itself. Specifically, the proposed rule sets a numerical threshold ("using more than 10^{26} computational operations") to require reporting on a dual-use foundation model training run, but the discussion of the proposed rule includes additional definitions and discussion of dual-use foundation models, including the definition in the Executive Order 14110 that potentially draws in a much broader set of models and associated training runs.¹

BIS should further clarify that this rule only applies to the original developers of large models. We have heard concerns about how this rule will be interpreted or applied outside the largest dedicated AI companies and whether other companies adapting, implementing, or building on top of these models would also need to report on training runs. We believe there are no reporting obligations on implementers, deployers, and users of AI models in this proposed rule, which is supported by BIS's stated expectation that 15 or fewer companies will exceed the reporting threshold.

BIS should implement the numerical thresholds for reporting but continue to explore alternatives to numerical reporting and risk thresholds

As ATAI noted in our April 29th comments to BIS, large models are not inherently riskier, and many narrowly scoped, highly capable models are more capable than general models of similar size.² Similarly, the size of a particular model training run is not likely to be a reliable predictor of

¹ "Executive Order on the Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence," 30 October 2023, <https://www.whitehouse.gov/briefing-room/presidential-actions/2023/10/30/executive-order-on-the-safe-secure-and-trustworthy-development-and-use-of-artificial-intelligence/>.

² "ATAI Submits Comments on Bureau of Industry and Security Proposed Rulemaking on Malicious Cyber-Related Activities," 29 April 2024, <https://alliancefortrustinai.org/atai-bis-rulemaking-comments/>.

risk. In those comments, we emphasized that it is more important to consider deployment contexts when assessing risk. Like other technologies, AI models are tools, and the context in which the model is deployed will almost always play a role in the risk involved, often more than the model's size.

Many AI and risk experts and government officials have expressed similar concerns. National Telecommunications and Information Administration (NTIA) Administrator Alan Davidson raised doubts about the numerical threshold as part of the definition of a “dual use foundation model” in the Biden administration's AI Executive Order.³ Davidson criticized the focus on how many parameters are used to train a model and recommended a more basic inquiry to understand AI risk. He did not suggest an alternate threshold.

We recognize that a numerical threshold is a reasonable means to limit the rule's scope at this time. Most other thresholds, such as the scope of the definition of a “dual-use foundation model,” would likely be too ambiguous to be effective and introduce additional confusion. However, we recommend exploring alternatives to a numerical threshold because risk and model size are not inherently connected.

Reporting Burdens

BIS should consider ways to reduce reporting burdens under the proposed rule

This NPRM will place onerous burdens for reporting on smaller companies developing AI models, particularly for those that do not have the resources to submit quarterly reports, which will be incredibly resource-intensive. Additionally, we believe that the estimated hourly burden that BIS puts forth vastly underestimates the actual burden of quarterly reporting. ATAI instead recommends annual reporting requirements.

These burdens may create a significant disincentive for innovation as companies evaluate the reporting burdens and choose not to undertake covered training runs. This reduction in research and innovation will have downstream effects, particularly for entities that do not develop their models but instead rely on implementing and iterating others' models, leading to an overall loss in momentum for AI innovation and development.

While BIS currently does not expect the requirements to apply to more than 15 companies, as AI ecosystem participation increases, more companies will likely fall under these requirements. If the burden of reporting outweighs the gains for a company, they may decide to focus innovation and development elsewhere and away from large AI models.

³ “NTIA's Davidson: Executive order errs in use of numeric thresholds to manage AI risk,” *Inside AI Policy*, 13 September 2024, <https://insideaipolicy.com/share/16863/>.

ATAI strongly recommends that BIS consider the impacts that these reporting burdens may have. There is a better balance for these reporting requirements that will still yield valuable information while allowing increased innovation and development.

Conclusion

Thank you for the opportunity to comment on these questions. If you have questions or believe that we can be helpful to your work in any way, please contact ATAI's coordinator, Heather West, at hewest@venable.com.