Answering 6 questions + Closing statement

2.i. The biggest risk in AI right now is from corporations. With significant resources that far outstrip what is available to the open weight models, corporations can create AI models that are truly dangerous in a way that open weight models are not. As an individual, I can afford to create a system that can at most train a model (currently) that is comparable to GPT 3.5, while organizations like OpenAI, Google, Microsoft, Apple, and X could already be sitting on AGI class AI's that they simply have not told us about. The profit motive of these organizations creates perverse incentives that lead them to create products that focus on profit before anything else, which can have a terrible effect on society. To put it bluntly, these organizations are effectively 'paperclip maximizers' for money and having closed systems that are significantly smarter and more powerful that they can use internally for their money maximization motives could lead these orgs to do things like affect politics domestically and abroad through subtly opinionated AI models in conjunction with more traditional PR and marketing techniques, manipulate markets on a grand scale, destroy whole sectors of the economy, etc. Any risk that an open model exposes is significantly worse in a closed model.

3.a. Because of resource constraints on the actors working on open models vs closed models, it forces them to explore more efficient ways of training and serving models that have benefits to all players in the space.

Take this development that Answer.AI released.

https://www.answer.ai/posts/2024-03-06-fsdp-qlora.html

This sort of innovation that allows users with 48GB of VRAM train 70B parameter models by combining FSDP, a method developed by Meta, and QLoRA, developed by Tim Dettmers et. All, is something that would not have happened outside of the context of open models due to a variety of reasons that they outline in the link.

3.b.c.d.e. Making model weight widely available derisks AI as it provides a counterweight to the risk that closed source AI poses to society while fostering actually open ai research that can be verified and built upon by experts worldwide. c) With the ability to retrain and finetune open models, corporations and individuals can tailor their AI systems to actually do what they intend it to do, rather than be forced to use an off the shelf product that may or may not be built for that purpose. d) By keeping things in the open, it could potentially accelerate AI development by unveiling techniques such as the one I mentioned in 3.a. This supports the United States' national security because in the AI race, whoever gets to AGI first and is able to deploy the model at scale will become the most powerful actor, essentially overnight. By having this sort of development be open, and in the current environment where the United States has access to the best hardware, it increases the likelihood that the United States will reach this point before actors in other countries. e) When training data and source code is available, it makes the development of AI more collaborative and allows smaller scale entities universities and smaller corporations the ability to build and iterate on models in the open.

Closing Statement:
It is imperative that as AI systems get more and more powerful and approach AGI, that further development and large scale training of these systems happens out in the open and with actual AI risk experts helping to direct the development of these systems. Researchers like Eliezer Yudkowsky should be heeded when they say that AI poses an existential risk for humanity. AI researchers and executives who work for corporations cannot be trusted to put proper safety measures in check because their incentives run counter to the incentives of the rest of society. Because that ship has sailed, all that we have left are open models to act as an extremely weak counterweight to closed models.