Comments of the Recording Industry Association of America on the National Telecommunications and Information Administration Request for Comment on Dual Use Foundation Artificial Intelligence Models with Widely Available Model Weights

Docket No. NTIA–2023–0009

Delivered via regulations.gov

March 27, 2024

The Recording Industry Association of America (RIAA) welcomes this opportunity to provide comments to the National Telecommunications and Information Administration (NTIA) request for comment on dual use foundation Artificial Intelligence (AI) models with widely available model weights.

RIAA is the trade organization that supports and promotes the creative and commercial vitality of music labels in the United States, the most vibrant recorded music community in the world. Our membership – which includes several hundred companies, ranging from small-to-medium-sized enterprises to global businesses – creates, manufactures, and/or distributes sound recordings representing the majority of all legitimate recorded music consumption in the United States. In support of its mission, the RIAA works to protect the intellectual property and First Amendment rights of artists and music labels; conducts consumer, industry, and technical research; and monitors and reviews state and federal laws, regulations, and policies.

I. Introduction

Human creative expression is at the core of what our members do and support and is vital to every nation's culture and economy. Just in the United States, the total copyright industries, including the creative sector, accounted for 8.14% of all U.S. employment in 2021, with wages exceeding the U.S. average annual wage by around 29%.[1] Part of the reason for this economic boost is our nation's strong protection of intellectual property rights and other rights of creators.

AI has the potential both to significantly increase creators' ability to express themselves artistically and to significantly harm them and rightsholders through the unethical development and deployment of AI systems that unfairly exploit their work, expression or identity.[2]

[1] Robert Stoner et al., *IIPA, Copyright Industries in the U.S. Economy, 2022 Report*, Secretariat Economists, prepared for the International Intellectual Property Alliance, p. 8, Dec. 2022, https://www.iipa.org/files/uploads/2022/12/IIPA-Report-2022_Interactive_12-12-2022-1.pdf.

[2] We have already experienced some of these benefits and harms firsthand. Please see the following comments the creative sector has made with National Telecommunications and Information Administration ("NTIA"), Office of Science Technology and Policy ("OSTP"), and Copyright Office for a description of the benefits and harms we have experienced to date: (a) June 12, 2023 comments in response to NTIA AI Accountability Policy Request for Comment, Docket No. NTIA-2023-0005, available at https://www.regulations.gov/comment/NTIA-2023-0005-1277; (b) July 7, 2023 comments in response to the OSTP Request for Information on National Priorities for Artificial Intelligence, Docket No. OSTP-TECH-2023-0007, available at https://www.regulations.gov/comment/OSTP-TECH-2023-0007-0231; and (c) the Oct. 30, 2023 and Dec. 6, 2023 comments to the Copyright Office in the Matter of Artificial Intelligence and Copyright, available at

Accordingly, it is important that creators and those that support them are at the table during any AI policy discussions. With this in mind, we offer the following comments.

II. Core Principles

Consistent with the Administration's views, NTIA should ensure that development and deployment of dual use foundation models is done "safely and responsibly,"[3] in a manner that "advance[s] human rights and human dignity,"[4] and does "not come at the price of civil rights or democratic values."[5] This necessarily includes respecting the rights of creators, performers, and other rightsholders in their creations, identities, and dignity.[6, 7] The G7 has espoused similar views in the AI development and deployment context, calling for "appropriate data input measures and protections for personal data and intellectual property,"[8] "[a]ppropriate transparency of training datasets,"[9] and respect for "rights related to privacy and intellectual property, including copyright-protected content."[10] The United Nations also adopted similar views in a resolution on AI, calling for the promotion of safe, secure and trustworthy AI systems that respect human rights and encourage, where appropriate and relevant, the implementation

---

https://www.regulations.gov/comment/COLC-2023-0006-8833 and https://www.regulations.gov/comment/COLC-2023-0006-8833 ("Copyright Office Comments").

[3] Exec. Order No. 14110, 88 Fed. Reg. 75191 (Nov. 1, 2023) (hereinafter "Executive Order").

[4] *See* Remarks by Vice President Harris on the Future of Artificial Intelligence, London, United Kingdom, Nov. 1, 2023, https://www.whitehouse.gov/briefing-room/speeches-remarks/2023/11/01/remarks-by-vice-president-harris-on-the-future-of-artificial-intelligence-london-united-kingdom/.

[5] *See* The White House, Office of Science and Technology Policy, *Blueprint for an AI Bill of Rights*, Oct. 4, 2022, available at https://www.whitehouse.gov/ostp/ai-bill-of-rights/. *See also* the statement from NTIA that it desires to develop AI accountability policies that can demonstrate "that AI systems are legal, effective, safe and otherwise trustworthy." 88 Fed. Reg. 22433 (Apr. 13, 2023).

[6] Those rights are also enshrined in the U.N. Universal Declaration of Human Rights. *See* U.N. Universal Declaration of Human Rights, art. 27, § 2 ("Everyone has the right to the protection of the moral and material interests resulting from any scientific, literary or artistic production of which he is the author"), art. 12 ("No one shall be subjected to arbitrary interference with his privacy, family, home or correspondence, nor to attacks upon his honour and reputation"); https://www.un.org/en/about-us/universal-declaration-of-human-rights. In addition, rights in creative expression are protected by the U.S. Constitution and federal law. *See* U.S. Const. art. 1, § 8, cl. 8 ("To promote the Progress of Science and useful Arts, by securing for limited Times to Authors and Inventors the exclusive Right to their respective Writings and Discoveries"), and Title 17, United States Code.

[7] We appreciate that NIST has already acknowledged the importance of intellectual property, including name, image, voice, and likeness rights, and the need to consider the risks associated with infringement of those rights in its AI Risk Management Framework ("AI RMF"). *See* AI RMF, p. 16, ("Training data may also be subject to copyright and should follow applicable intellectual property rights laws") and p. 24 (includes in the govern function a recommendation that "[p]olicies and procedures are in place that address AI risks associated with third-party entities, including risks of infringement of a third-party's intellectual property or other rights"), available at https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf.

[8] *See Hiroshima Process International Guiding Principles for Organizations Developing Advanced AI system*, ¶ 11, Oct. 30, 2023, available via https://digital-strategy.ec.europa.eu/en/library/hiroshima-process-international-guiding-principles-advanced-ai-system.

[9] *Id*.

[10] *See Hiroshima Process International Code of Conduct for Organizations Developing Advanced AI Systems*, ¶ 11, Oct. 30, 2023, available via https://digital-strategy.ec.europa.eu/en/library/hiroshima-process-international-code-conduct-advanced-ai-systems.

of appropriate safeguards to respect intellectual property rights, including copyright protected content.[11]

Fundamentally, responsible AI innovation and deployment requires a human-centric approach. As a core principle, AI policy should encourage AI developers and deployers to create AI tools that *support* human creativity and should discourage those that attempt to *supplant* human creativity or exploit human creativity without authorization.

That is why we, along with more than 180 other entities, have signed onto the Human Artistry Campaign.[12] The Human Artistry Campaign calls for policy makers, AI developers, and those that deploy AI to take into account the following principles:

(i) technology has long empowered human expression, and AI will be no different;
(ii) human created works will continue to play an essential role in our lives;
(iii) use of copyrighted works and the use of voices and likenesses of professional performers requires authorization and free-market licensing from all rightsholders;
(iv) governments should not create new copyright or other IP exemptions that allow AI developers to exploit creations without permission or compensation;
(v) copyright should only protect the unique value of human intellectual creativity;
(vi) trustworthiness and transparency are essential to the success of AI and protection of creators; and
(vii) creators' interests must be represented in policy making.[13]

With those principles in mind, we offer the following more specific comments.

III. Comments

The following comments are responsive to questions 2 and 3 of the request for comment.[14]

As an initial matter, we reiterate the fundamental importance of transparency – AI developers and deployers must be required to collect, maintain, and share (as appropriate) proper records concerning the entire AI development and deployment lifecycle.[15] This includes not only information about the model weights/parameters, but also complete documentation of:

[11] "General Assembly Adopts Landmark Resolution on Artificial Intelligence", United Nations, March 21, 2024, https://news.un.org/en/story/2024/03/1147831 (the text of the resolution is available at https://documents.un.org/doc/undoc/ltd/n24/065/92/pdf/n2406592.pdf?token=Spve3fWXZ9HUT97mCU&fe=true).
[12] *See* https://www.humanartistrycampaign.com/.
[13] *Id.*
[14] Question 2 asks "How do the risks associated with making model weights widely available compare to the risks associated with non-public model weights." Question 3 asks "What are the benefits of foundation models with model weights that are widely available as compared to fully closed models."
[15] See comments identified in footnote 2 above.

(i) what materials were ingested to develop the AI system (or to fine tune or adapt a pre-trained AI system) and in what manner, [16]
(ii) the provenance of such materials, including whether any licenses or authorizations were sought or obtained to authorize such use and copies of those licenses or authorizations,
(iii) the articulated rationale for selecting and using the materials ingested for the AI system's development,
(iv) the articulated purpose of the AI model itself and its intended outputs,
(v) the AI system's overall functioning, including potentially infringing functionality,
(vi) the individual or organization responsible for the AI system (including who is responsible for ingesting the materials, who is responsible for any foundational AI model, who is responsible for any fine tuning of the AI model, who is deploying the AI system, etc.),
(vii) risk assessments concerning the potential misuse and abuse of such a model,
(viii) what processes were used, and what decisions were made, during the AI system development and deployment, and
(ix) such other information relevant to explain, interpret, and understand the decision-making or generative process used within the AI system, and its resulting outputs, in an explainable fashion[17] to make it reproduceable, tractable, and ultimately trustworthy.

Maintaining such records should be a requirement, regardless of whether the weights are made widely available or not. This helps not only with auditing the provenance of training materials and whether such materials were appropriately licensed, but also with evaluating and managing other risks associated with the model, such as bias, privacy risks, errors, etc.

In addition, to further promote transparency, AI developers and deployers of dual use foundation models should, at a minimum, disclose to the public and regulators the purpose of the AI system and its overall functionality (including its potential to violate third-party intellectual property, privacy or other third-party rights), who is the individual or entity responsible for the AI system and their location and contact information, the provenance of the materials ingested during the AI system's development, and basic information needed to provide algorithmic transparency. This information should be disclosed to regulators, interested parties, and the general public at a level of specificity necessary to address their legitimate concerns relating to the development, training, and operation of AI systems, including to help in assessing whether or not the weights should be made widely available. AI developers and deployers should also make more detailed records and audits of the AI system available to third parties who have a good faith claim that their rights have been violated in

[16] This information should be granular enough that a rights owner can review the records and determine if their works have been ingested and the extent to which they have been relied on in generating a particular output. For recorded music, that means that AI developers would have to track fields including, but not necessarily limited to, track title, artist name, songwriter name, International Standard Recording Code ("ISRC"), International Standard Works Code ("ISWC"), P-line (identifying the copyright owner of a sound recording), and C-line (identifying the copyright owner of the underlying musical composition).
[17] This is often referred to as explainable artificial intelligence or XAI.

connection with the AI system's development or deployment,[18] and any such violations should be taken into account in considering whether the weights should be made widely available.

In considering whether (and under what circumstances) to make the foundation model weights widely available, there are clear risks and benefits that must be factored into any such decision. Below we highlight some of those risks and benefits as they relate to models that can be used to generate music or audiovisual works that include music:

- To ensure that our valuable copyrighted works are not used by AI developers and deployers without authorization and without appropriate compensation, the music community is concerned, first and foremost, with knowing exactly which copyrighted works have been used to develop or fine tune a given model, the provenance of those materials, whether those materials were properly licensed or cleared for such use, and how those contributions to the model are weighted. In our experience to date, generative AI developers rarely disclose any of that information. This problem could be partially solved by implementing recordkeeping and transparency requirements that obligate AI developers to disclose such information to appropriate third parties.
- Although we recognize that making model weights widely available can be seen as a step towards greater transparency, unless the AI developers and deployers have secured appropriate and enforceable licenses permitting them to ingest and otherwise use copyrighted music or music videos in their AI systems, making the weights widely available (and the accompanying code) is likely to exponentially exacerbate the harms to creators and rightsholders and their ability to respond to protect their rights. Widely available access to weights (and the accompanying code) would make it easier for subsequent developers and deployers to develop or fine tune models that further infringe or act as substitutes for those copyrighted works by permitting them to modify, fine tune, or deploy AI tools that piggyback off the underlying unlicensed dataset. This problem becomes more acute when one considers that a rightsholder might need to seek enforcement in jurisdictions outside the United States to protect their rights, and that the number of subsequent developers and deployers may be so vast that enforcement becomes practically impossible. On the plus side, making such weights widely available for bona fide research purposes could help researchers, such as copyright owners, determine which works were used to develop a given model, evaluate what is their relative contribution to the model, and better assess other risks and benefits posed by the model.
- Where a model is developed using *licensed* copyrighted music or music videos and the licenses specify the terms and conditions under which the weights can be made widely available, the risks noted above are mitigated to a significant extent. In such a case, the license terms would presumably flow through to whomever uses the weights (and associated code) to copy or fine tune the model. This would benefit the rightsholders to help ensure that they are appropriately compensated for the use of their works, those using the widely available weights to copy or fine tune the model to reduce the time and resources needed for further development, and model users by creating a more

[18] *See* Copyright Office Comments for a proposal for a third-party subpoena process to obtain such information.

competitive environment for the use of such models. If the license terms do not address downstream uses or the terms concerning downstream uses are disregarded, then the risks noted above with respect to widely available weights would apply as those using the weights could bypass licensing the works for AI development without any compensation to the rightsholders.

- The same analysis would apply to the unauthorized use of an artist's name, image, likeness, or voice. Where those rights are not cleared for the dual use foundation model, making the model weights widely available would exacerbate those harms.[19] On the other hand, where those rights have been appropriately cleared, including for use where the weights are made widely available, making the weights (and associated code) widely available should have a positive benefit.
- There is one additional risk associated with the wide availability of model weights. In an effort to better control the potential outputs of AI models, dual use AI foundation model developers have recently begun placing guardrails around either the prompts they will accept as inputs to their models or on the generated output that will be revealed to the user. Making the weights widely available can make it easier for an unscrupulous developer to modify or fine tune the model in a manner that bypasses those guardrails.

* * *

We thank NTIA for the opportunity to respond to this RFC. Please let us know if you have any questions or if we can provide any additional input. We look forward to continuing this conversation with NTIA and other policy makers as AI technology and its impact on the marketplace and society evolves.

Respectfully submitted,

/Victoria Sheckler/
Victoria Sheckler
SVP, Deputy General Counsel
Recording Industry Association of America

[19] These harms include expanded availability of models to create unauthorized deep fakes, disinformation, cybersecurity vulnerabilities through the unauthorized use of a celebrity or political figure's persona, etc.